Does letting the model think for longer remove the conflict between its instructions?

Every prompt asks for exactly 4 sentences, plus a competing instruction: make each sentence longer than the last. The two lines differ only in one further formatting instruction.

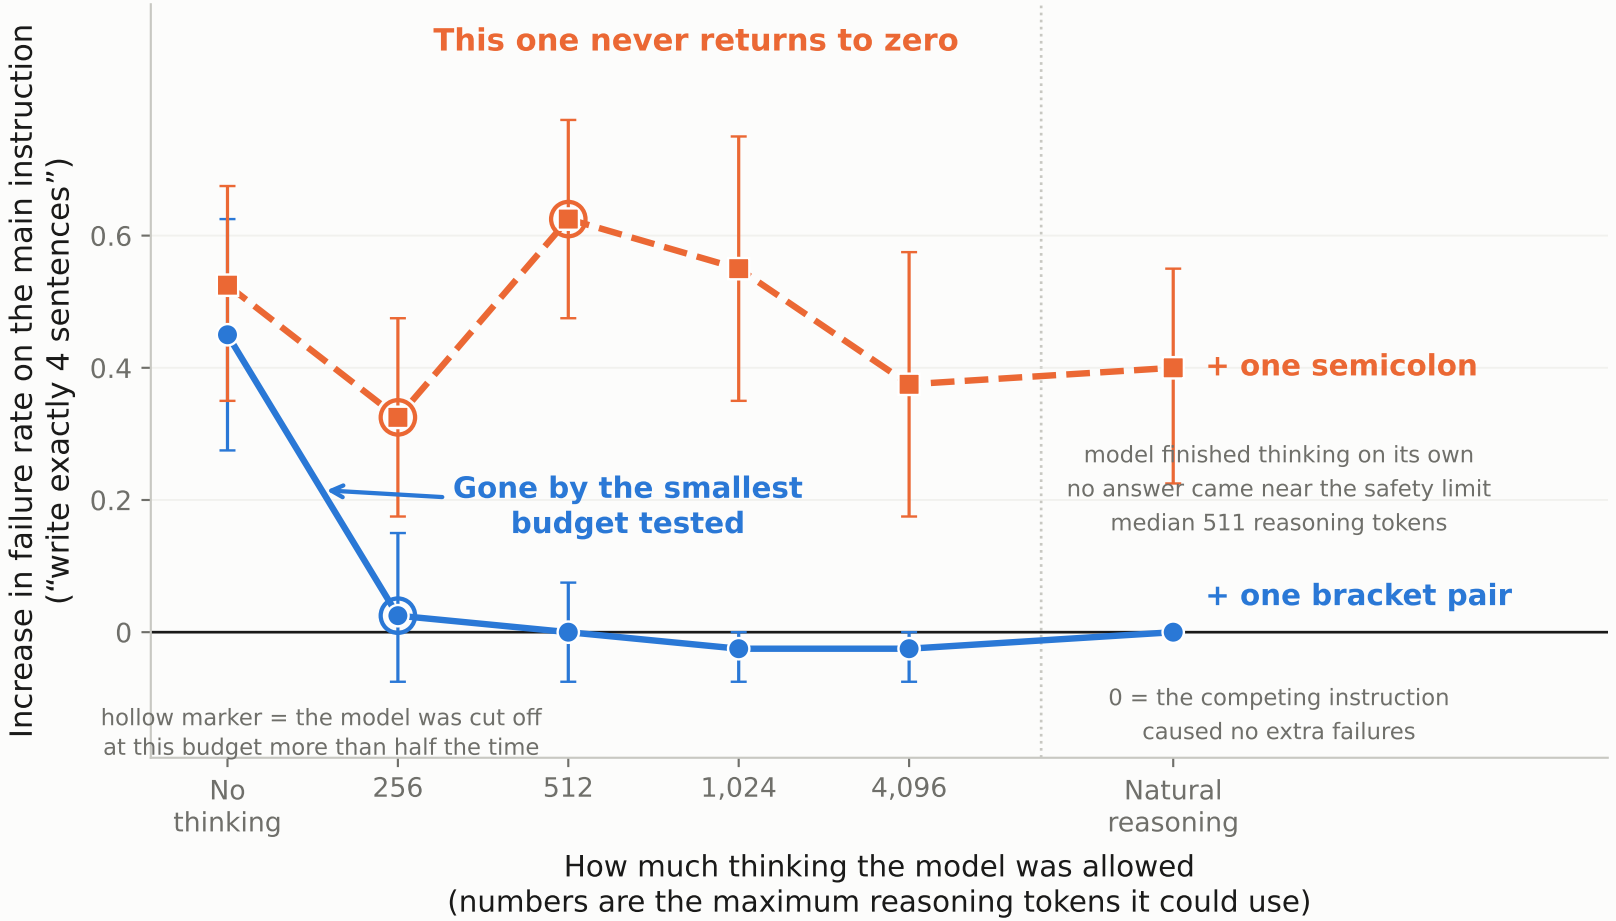

Qwen3-4B, one instruction pair · every point is a paired difference across 20 tasks (2 answers per task per condition, 80 answers per point) · error bars are 95% intervals from resampling the tasks
Budgets actually run: 0, 256, 512, 1,024, 4,096 — 2,048 was not run, and the axis is ordered, not to scale