Does letting the model think for longer remove the conflict between its instructions?

Every prompt asks for exactly 4 sentences, plus a competing instruction, plus one formatting instruction. The two lines differ only in that formatting instruction — and in nothing else.

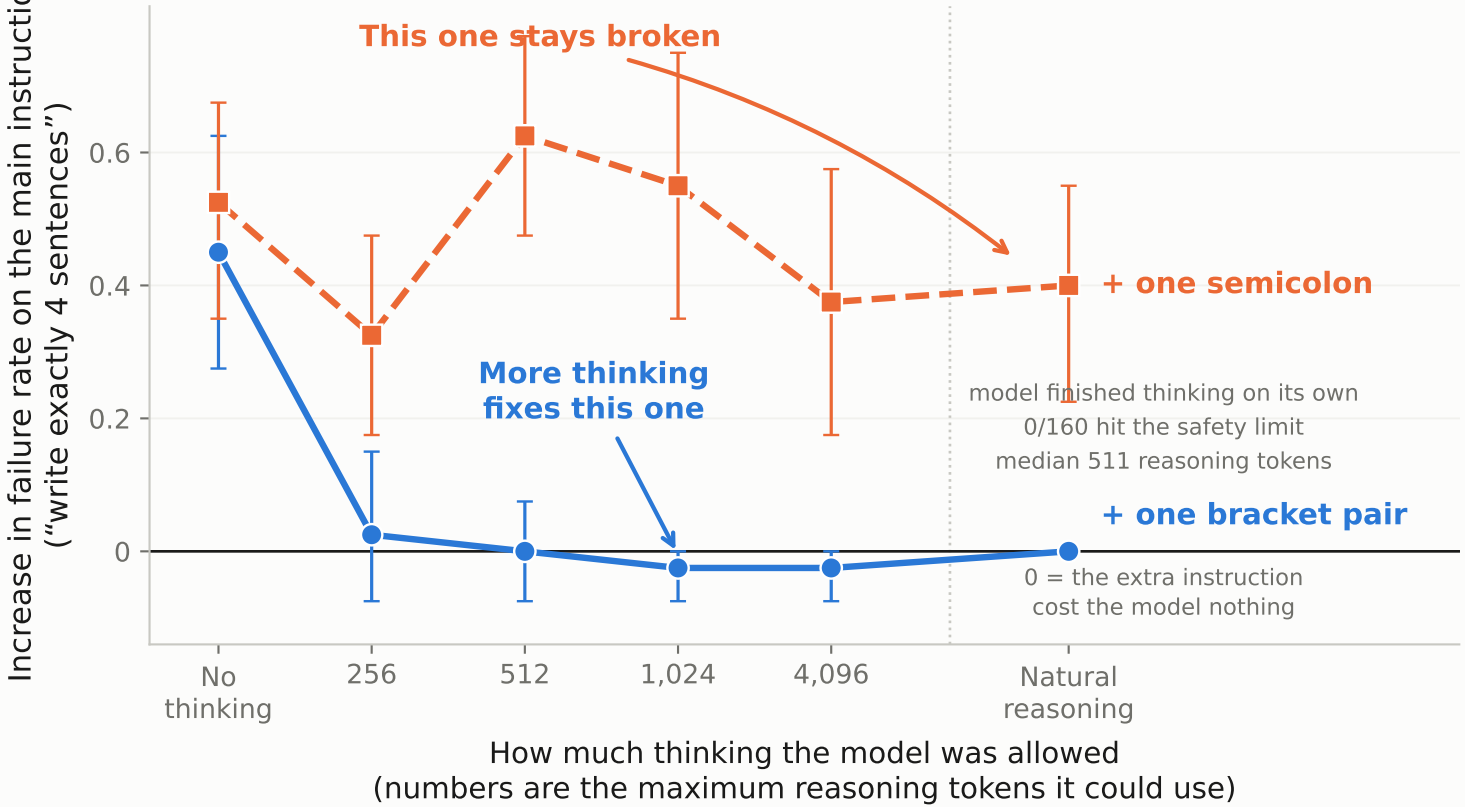

Qwen3-4B · 20 tasks × 2 replicates behind every point · error bars are 95% intervals from resampling the tasks