Additional test-time reasoning resolves ORDINARY directional interference but not AMPLIFIED interference
Panels C and D are the result; A, B, E and F show what it is made of

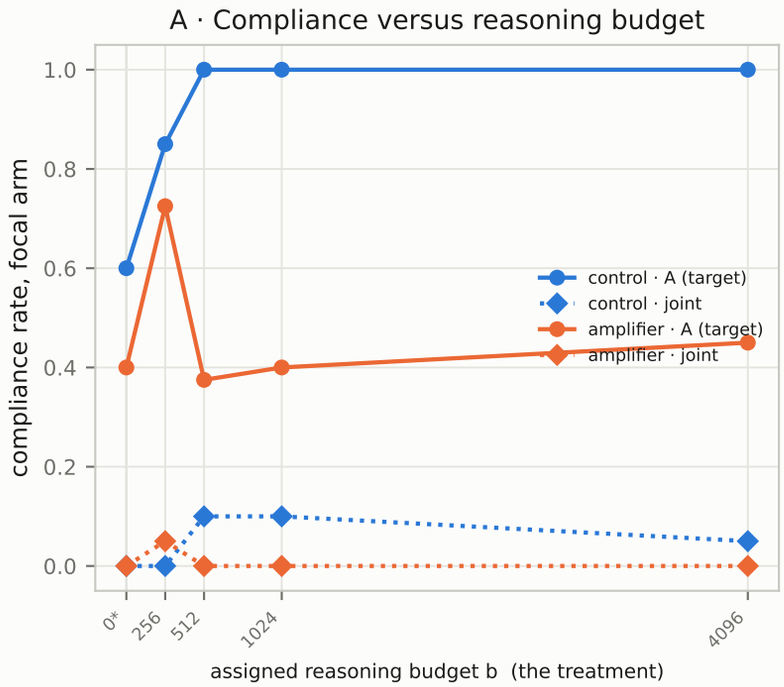

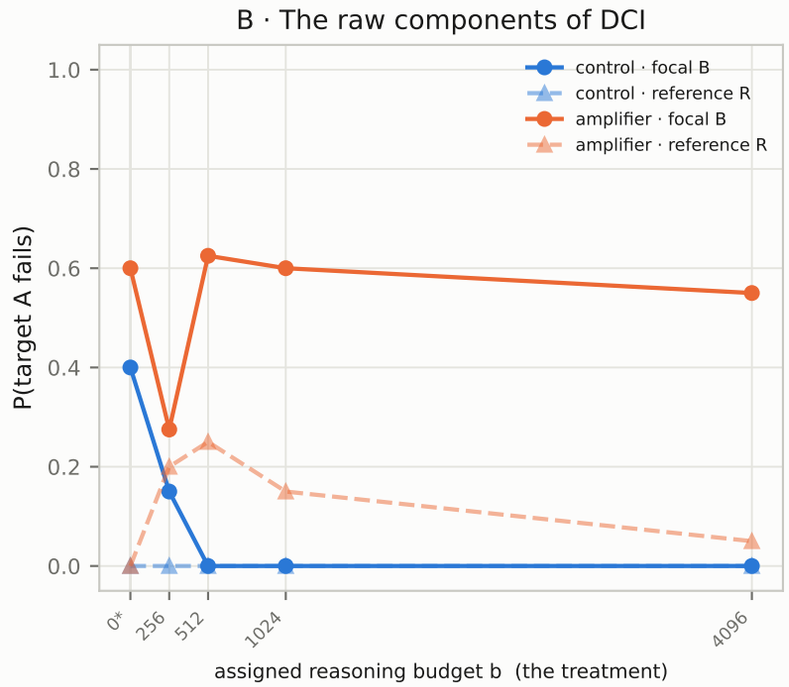

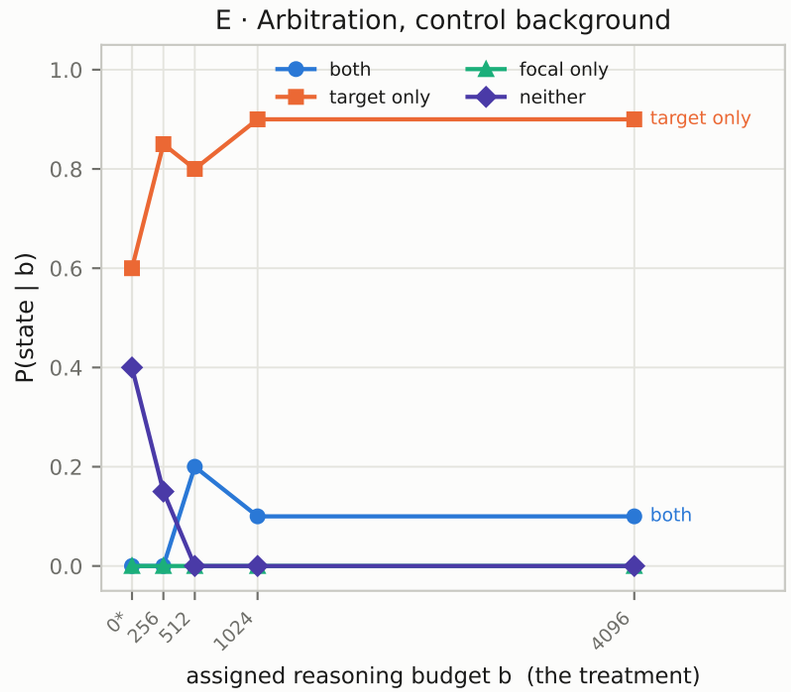

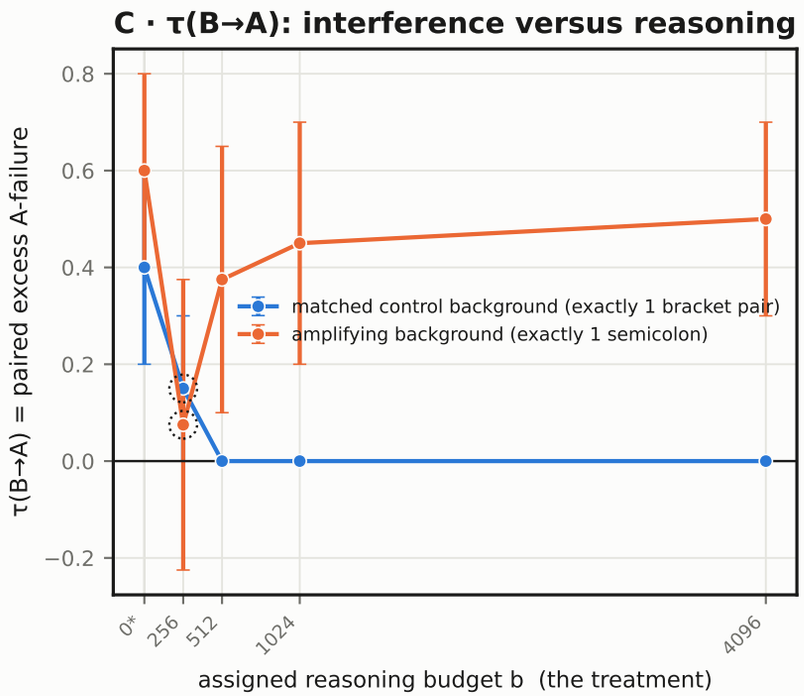

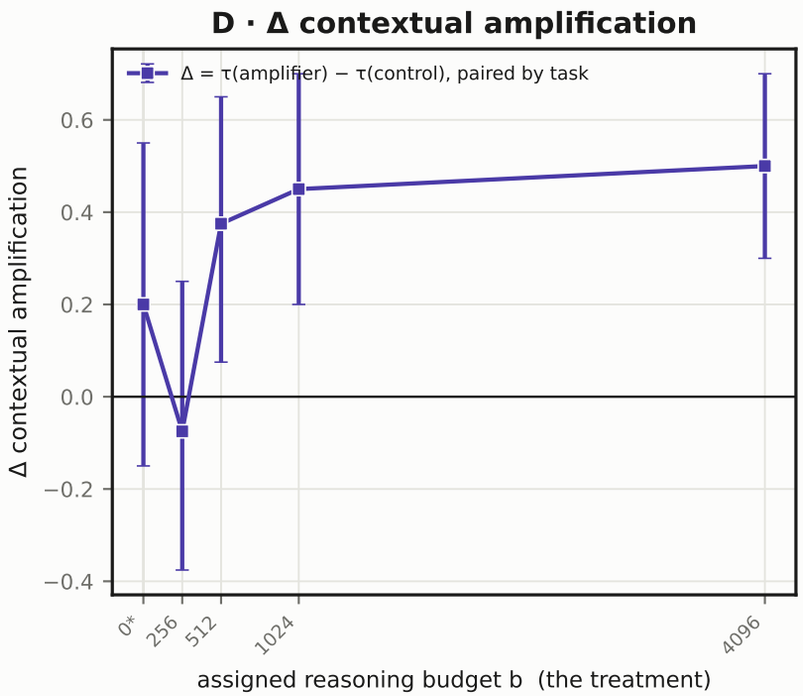

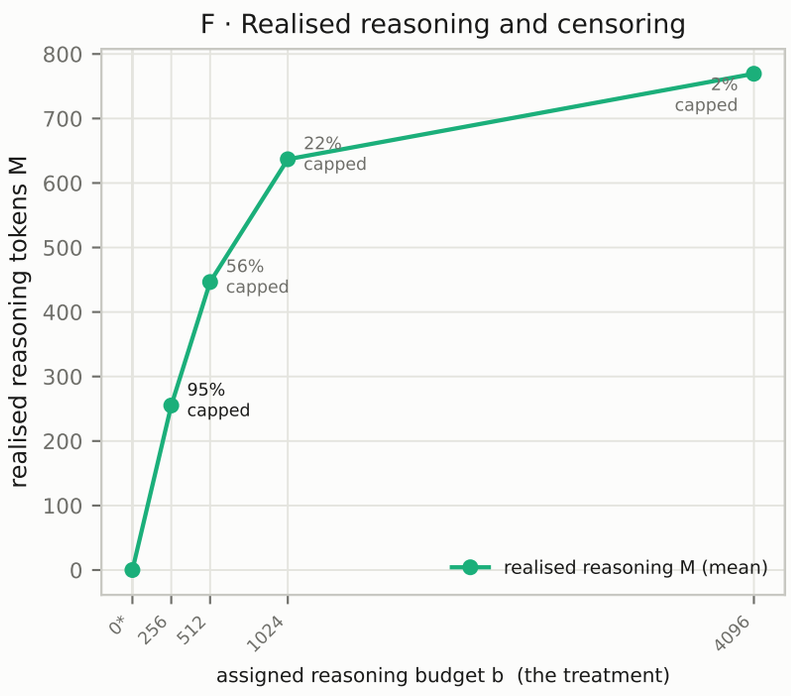

* b = 0 is the model's native non-thinking mode, not a zero-length truncation

Reasoning panels: Qwen3-4B @ 1cfa9a72, vLLM 0.15.1+rocm720, tensor-parallel 2 (disable_custom_all_reduce, NCCL_P2P_DISABLE=1), bf16, T=0.7, top-p 0.9, fixed 512-token answer budget · b=0 is the model's native non-thinking mode

Pair 3: A = exactly 4 sentences (target) · B = strictly increasing sentence lengths (focal) · C_L3 = matched reference · π = QAC, k = 3 · 20 pilot tasks × 2 replicates · intervals: percentile bootstrap over 20 task clusters, 4000 draws — exploratory, uncorrected