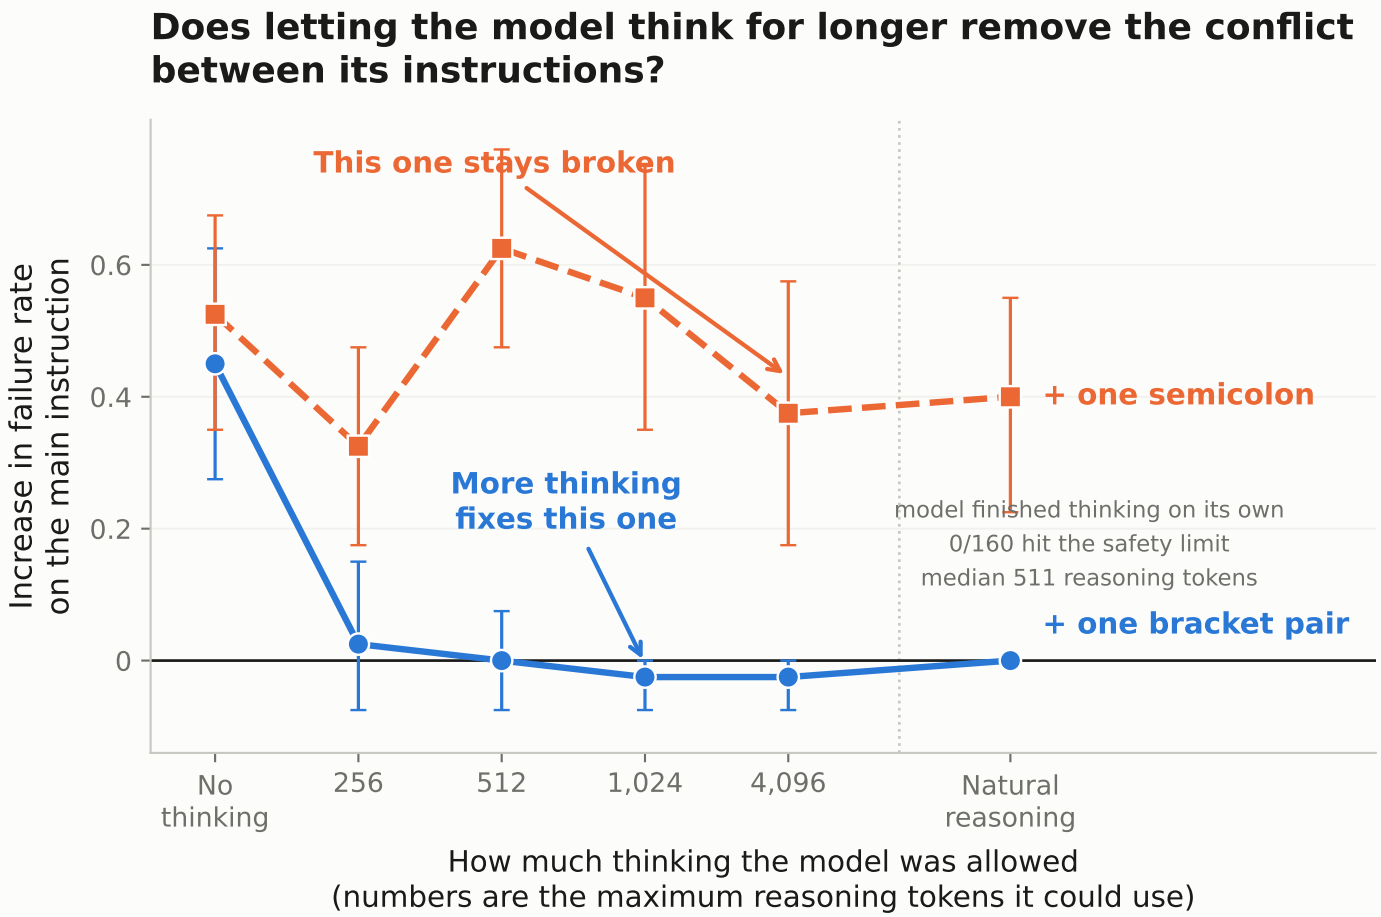
Does letting the model think for longer remove the conflict between its instructions?
This one stays broken
+ one semicolon
More thinking fixes this one
model finished thinking on its own
0/160 hit the safety limit
median 511 reasoning tokens
+ one bracket pair
Increase in failure rate on the main instruction
0.6
0.4
0.2
0
No thinking
256
512
1,024
4,096
Natural reasoning
How much thinking the model was allowed
(numbers are the maximum reasoning tokens it could use)