Does letting the model think for longer remove the conflict between its instructions?

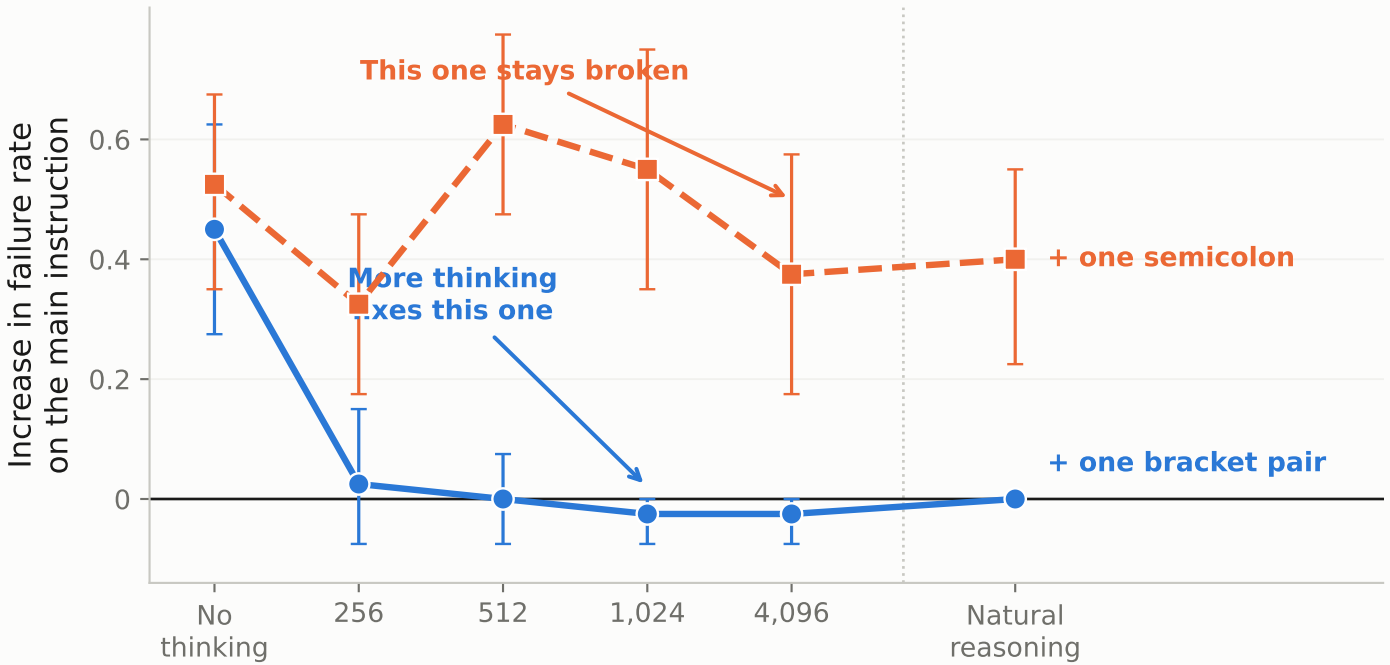

What the vertical axis measures

Prompt A
Write exactly 4 sentences.
Begin every sentence with a
word starting with S.

↓

Prompt B
Write exactly 4 sentences.
Make each sentence longer
than the one before.

How much more often does the model get "exactly 4 sentences" wrong in Prompt B than in Prompt A?

Both prompts also carry the same third instruction — either "one bracket pair" or "one semicolon".